

Alp Cetin

San Francisco

falpercetin@gmail.com · github.com/mottopanikeiku · linkedin.com/in/fatihctn

I'm building Tenth Man Labs. We put probabilities on geopolitical events for investment, insurance, and risk teams, with the sources behind each number and a dated record of how it moved.

My background is GPU systems. Lately most of my work is on evaluation and reinforcement learning.

	Tenth Man Labs tenthmanlabs.com Founder	2026–
	Most of the technical work is evaluation: scoring forecasts against questions that later resolve, and catching backtests that can see the future. In private beta, covering China and East Asia.	
Evals and RL	eval-power github.com/mottopanikeiku/eval-power	2026
	How many benchmark items it takes to tell two language models apart. Across 395 models on five benchmarks, almost no neighboring leaderboard gaps were significant, and studies planned from small pilots detected the difference only 49–79% of the time instead of 80%.	
	seed-power github.com/mottopanikeiku/seed-power	2026
	I tested RL study plans on 11,440 fresh training seeds. Plans targeting 80% power detected a difference in 67 of 119 executable comparisons (56.3%). A small pilot did not deliver the detection rate it promised.	
	control-clock github.com/mottopanikeiku/control-clock	2026
	Verge Lab github.com/mottopanikeiku/verge-lab	2026
	BranchPilot github.com/mottopanikeiku/branchpilot	2026
	Faultline github.com/mottopanikeiku/faultline	2026
	Quantile cycles github.com/mottopanikeiku/quantile-cycles	2026
Systems	cpu-decode github.com/mottopanikeiku/cpu-decode	2026
	A from-scratch C++ int8 decoder for Qwen2.5-0.5B on a laptop CPU. With 2 threads and short context it reaches 69 tokens per second, 83% of the memory-read ceiling and faster than llama.cpp; at 4,096 tokens of context llama.cpp wins, 44 to 32, because my attention is still scalar.	
	attention-numeric github.com/mottopanikeiku/attention-numeric	2026
	When FlashAttention-3's trick of rotating queries and keys before FP8 rounding makes attention worse. Across six open models it hurt 250 of 2,880 heads; a rounding-noise model committed in advance found the harmed heads in three unseen models with AUC 0.98. In Qwen2.5-0.5B, rotation raised perplexity 91% under emulated FP8 attention.	
	Aperture	2026
	Time to first token for a 27B model on one RTX 6000 Ada: under a second at p99 on cold 3,072-token prompts (803 ms), and 1.09 s at 4,096 tokens, which is where the claim stops.	
	HeliosTune github.com/mottopanikeiku/heliostune	2026
	Open source	2026
	Merged pull requests in Ray, Sentence Transformers, SGLang, and Celery.	
Research	Privacy-Preserving ML Lab Illinois Tech	2025–
	GPU and memory-system acceleration for secure inference with function secret sharing, where the keys for a transformer-scale model can run past 100 GB.	
	Computational Chemistry Lab Illinois Tech	2024–
	Molecular dynamics of the 5-HT _{2B} serotonin receptor for Alzheimer's drug discovery, and models that predict which signaling pathway a ligand will favor.	
Background	M.S., Artificial Intelligence Illinois Tech	2026
	B.S., Computer Science and Statistics Illinois Tech	2025
	Data science intern John Snow Labs	2023
	Taught programming to more than 50 students in six countries	2021–